



DPO beta study: 600 training pairs per condition

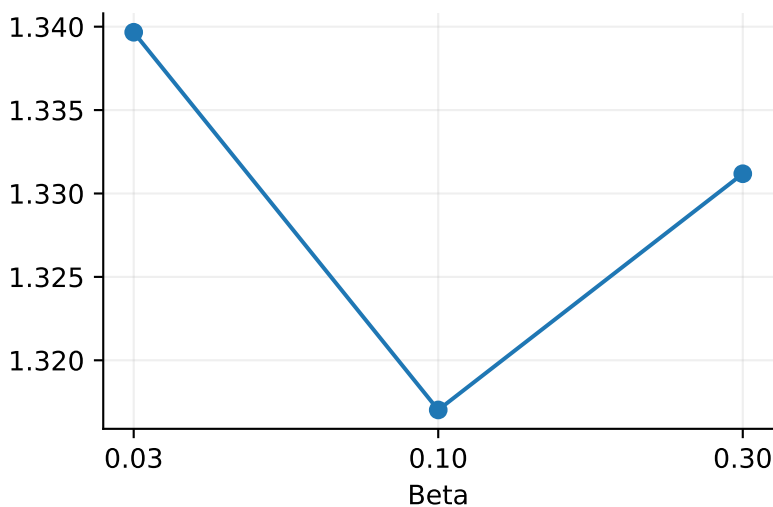
Preference accuracy (%)



Sampled KL estimate / token



Mean reward-model score



Mean response length (tokens)

