



Standard PPO: 20 updates; 0/40 policy and 2/40 critic steps skipped


Raw reward




Effective reward




Sampled KL diagnostic




Token entropy




Policy loss (epoch mean)




Critic MSE (epoch mean)



Online clip fraction



Response length (tokens)



Gradient norms before clipping

