




Raw reward



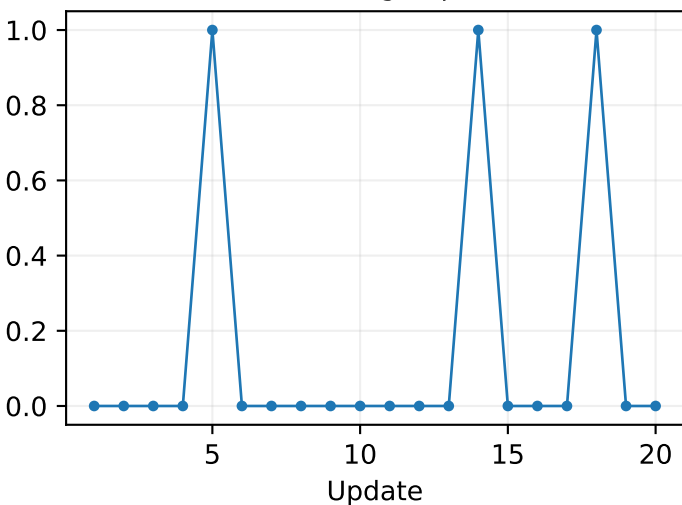
Raw sampled log-probability difference



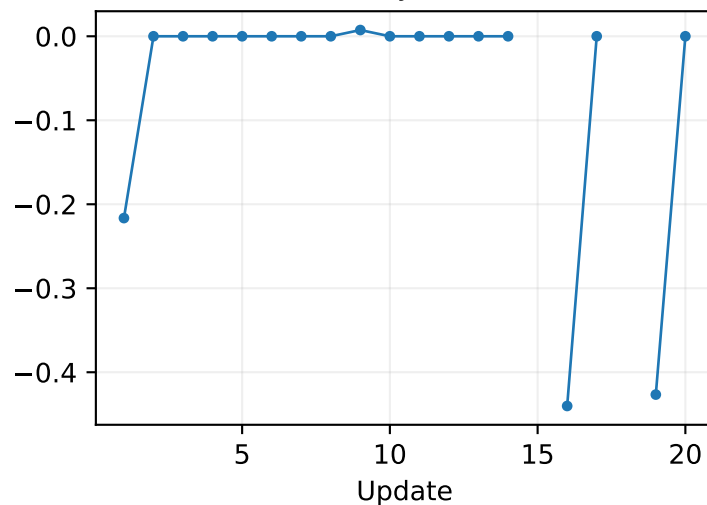
Within-group reward standard deviation



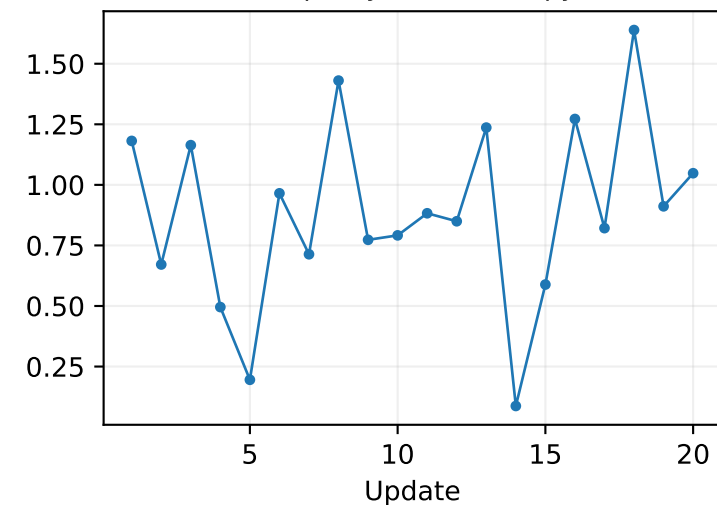
Uninformative-group fraction



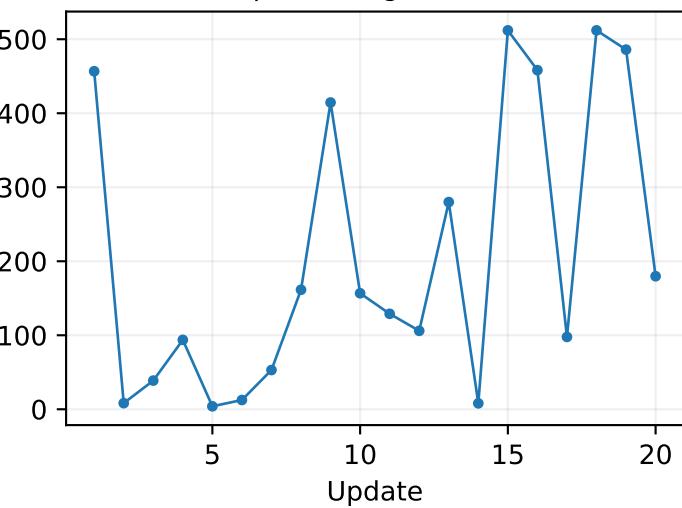
Policy loss



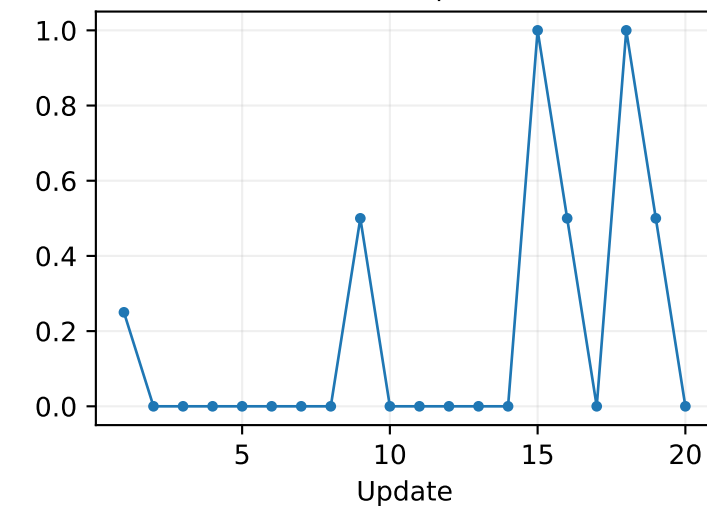
Full-policy token entropy



Response length (tokens)



Generation-cap fraction



Policy gradient norm before clipping

