



Fixed AI judge outputs: 450 prompts per policy (250 SAFE, 200 UNSAFE)
Class-inconsistent labels are retained; zero over-refusal is not a validated absence of over-refusal.


Safe answer | SAFE




Over-refusal | SAFE




Unsafe compliance | UNSAFE



Justified refusal | UNSAFE



Ambiguous judge labels | all



Mean response tokens

